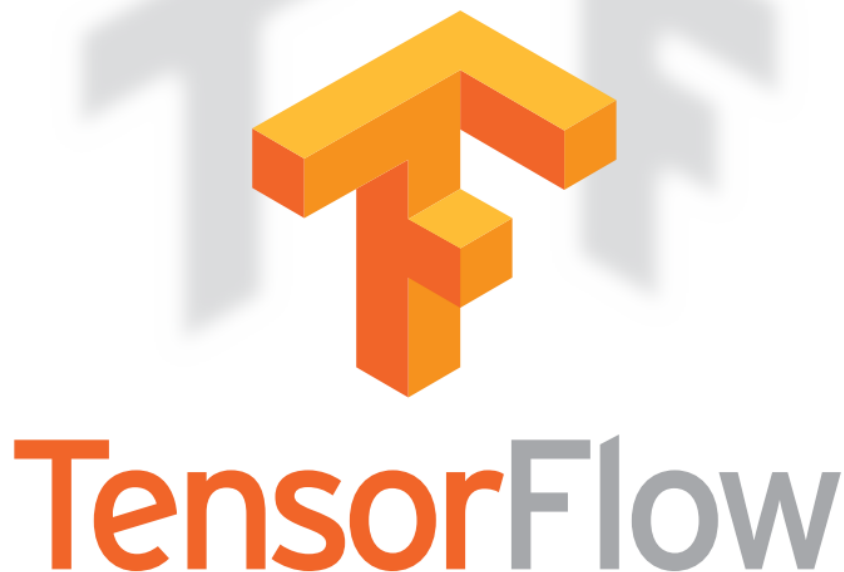


딥러닝 구현을 위한 텐서플로우 개발

김성균



딥러닝 구현을 위한 **텐서플로우** 개발

5강

머신러닝 실용과 MNIST 데이터셋


01.Learning rate

```
LEARNING_RATE = 0.01
```

```
# Cross entropy cost/loss
```


```
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
```

```
optimizer = tf.train.GradientDescentOptimizer(learning_rate = LEARNING_RATE).minimize(cost)
```




02. Large learning rate: overshooting

- cost가 증가하는 overshooting 현상



03. Small learning rate: takes too long, stops at local minimum




- 너무 조금씩 감소하는 현상



04. Try several learning rates


- cost function 관찰
- 경사도 하강속도 확인
- 다양한 learning rate을 사용해서 여러 번에 걸쳐 실행하는 것이 최선

05.경사도 하강을 위한 전처리(preprocessing)




05.경사도 하강을 위한 전처리(preprocessing)

x1	x2	Y
1	9000	A
2	-5000	A
4	-2000	B
6	8000	B
9	9000	C



05.경사도 하강을 위한 전처리(preprocessing)



06.표준화

- Normalization :

- 해당값을 0~1사이의 값으로 나타내는 척도법

$$x'_j = \frac{x_j - x_{min}}{x_{max} - x_{min}}$$

- Standardization:

- 속성값이 편견으로부터의 위치를 표준편차다위로 표현


$$x'_j = \frac{x_j - \mu}{\sigma}$$

μ_j : 평균
 σ_j : 표준편차

$$X_std[:,0] = (X[:,0] - X[:,0].mean()) / X[:,0].std()$$


07. Overfitting

- 모델이 학습데이터에 너무 잘 맞아서 학습데이터에만 최적화된 경우
- 정말 잘 맞추기 위해 과도하게 복잡해져서 실제로 사용할 때는 오히려 맞지 않는 현상




Under Fit

학습이 덜되어
error가 많이 발생함






Appropriate



Over Fit

학습데이터에 최적화
되어 error이 발생하지
않음

07. Overfitting




08. Overfitting의 해결방법

- 많은 학습데이터
- 입력으로 들어오는 변수(feature, x)의 갯수를 줄이기
- 정규화(Regularization)을 사용


09. 정규화(Regularization)

- W (weight)가 너무 큰 값들을 갖지 않도록 하는 것
- 비정상적이거나 애매한 위치에 있는 데이터를 올바르게 예측하는 것을 '맞았다'라고 얘기할 수는 없음
- 오히려 '틀렸다'라고 얘기하는 것이 더욱 좋을 수 있음
- 최소한의 에러를 인정하는 'Just right'



09. 정규화(Regularization)

- W(weight)가 너무 큰 값들을 갖지 않도록 하는 것
- 비정상적이거나 애매한 위치에 있는 데이터를 올바르게 예측하는 것을 '맞았다'라고 얘기할 수는 없음
- 오히려 '틀렸다'라고 얘기하는 것이 더욱 좋을 수 있음
- 최소한의 에러를 인정하는 'Just right'



- W(weight)가 너무 큰 값들을 갖지 않도록 하는 것
- cost 함수가 틀렸을 때 높은 비용이 발생할 수 있도록 벌점(penalty)을 부과

`L2reg = 0.001 * tf.reduce_sum(tf.square(W))`

$$Cost = \frac{1}{N} \sum_i \underbrace{D(S(WX_i + b), L_i)}_{\text{예측값}} + \lambda \sum W^2$$

실제값과의 차

X_i, L_i 입력값
 W, b 학습해야할 값

10. Training and Test datasets

- 학습용 데이터 세트와 테스트용 데이터 세트 구분

```
x_data = [[1, 2, 1], [1, 3, 2], [1, 3, 4], [1, 5, 5], [1, 7, 5], [1, 2, 5], [1, 6, 6], [1, 7, 7]]
y_data = [[0, 0, 1], [0, 0, 1], [0, 0, 1], [0, 1, 0], [0, 1, 0], [0, 1, 0], [1, 0, 0], [1, 0, 0]]
# Evaluation our model using this test dataset
x_test = [[2, 1, 1], [3, 1, 2], [3, 3, 4]]
y_test = [[0, 0, 1], [0, 0, 1], [0, 0, 1]]
```



```
X = tf.placeholder("float", [None, 3])
```

```
Y = tf.placeholder("float", [None, 3])
```

```
W = tf.Variable(tf.random_normal([3, 3]))
```

```
b = tf.Variable(tf.random_normal([3]))
```

```
# tf.nn.softmax computes softmax activations
```

```
# softmax = exp(logits) / reduce_sum(exp(logits), dim)
```

```
hypothesis = tf.nn.softmax(tf.matmul(X, W) + b)
```

```
# Cross entropy cost/loss
```

```
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
```

```
# Try to change learning_rate to small numbers
```

```
optimizer = tf.train.GradientDescentOptimizer(  
    learning_rate=LEARNING_RATE).minimize(cost)
```

```
# Correct prediction Test model
```

```
prediction = tf.argmax(hypothesis, 1)
```

```
is_correct = tf.equal(prediction, tf.argmax(Y, 1))
```

```
accuracy = tf.reduce_mean(tf.cast(is_correct, tf.float32))
```

```
# Launch graph
```

```
with tf.Session() as sess:
```

```
    # Initialize TensorFlow variables
```

```
    sess.run(tf.global_variables_initializer())
```

```
    for step in range(201):
```

```
        cost_val, W_val, _ = sess.run(  
            [cost, W, optimizer], feed_dict={X: x_data, Y: y_data})
```

```
        print(step, cost_val, W_val)
```

```
# predict
```

```
print("Prediction:", sess.run(prediction, feed_dict={X: x_test}))
```

```
# Calculate the accuracy
```

```
print("Accuracy: ", sess.run(accuracy, feed_dict={X: x_test, Y: y_test}))
```

```
199 0.672261 [[-1.15377033  0.28146935  1.13632679]
```

```
 [ 0.37484586  0.18958236  0.33544877]
```

```
 [-0.35609841 -0.43973011 -1.25604188]]
```

```
200 0.670909 [[-1.15885413  0.28058422  1.14229572]
```


```
 [ 0.37609792  0.19073224  0.33304682]
```

```
 [-0.35536593 -0.44033223 -1.2561723 ]]
```


```
Prediction: [2 2 2]
```

```
Accuracy: 1.0
```

11.Learning rate: NaN!



Large learning rate: Overshooting.



Small learning rate: Many iterations until convergence and trapping in local minima.

http://sebastianraschka.com/Articles/2015_singlelayer_neurons.html

```
X = tf.placeholder("float", [None, 3])
Y = tf.placeholder("float", [None, 3])
```

LEARNING_RATE = 1.5

```
W = tf.Variable(tf.random_normal([3, 3]))
b = tf.Variable(tf.random_normal([3]))
```

```
# tf.nn.softmax computes softmax activations
# softmax = exp(logits) / reduce_sum(exp(logits), dim)
hypothesis = tf.nn.softmax(tf.matmul(X, W) + b)
```

```
# Cross entropy cost/loss
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
# Try to change learning_rate to small numbers
optimizer = tf.train.GradientDescentOptimizer(
    learning_rate=LEARNING_RATE).minimize(cost)
```

```
# Correct prediction Test model
prediction = tf.argmax(hypothesis, 1)
is_correct = tf.equal(prediction, tf.argmax(Y, 1))
accuracy = tf.reduce_mean(tf.cast(is_correct, tf.float32))
```

```
# Launch graph
```

```
with tf.Session() as sess:
```

```
    # Initialize TensorFlow variables
```

```
    sess.run(tf.global_variables_initializer())
```

```
    for step in range(201):
```

```
        cost_val, W_val, _ = sess.run(
```

```
            [cost, W, optimizer], feed_dict={X: x_data, Y: y_data})
```

```
        print(step, cost_val, W_val)
```

```
    # predict
```

```
    print("Prediction:", sess.run(prediction, feed_dict={X: x_test}))
```

```
    # Calculate the accuracy
```

```
    print("Accuracy: ", sess.run(accuracy, feed_dict={X: x_test, Y: y_test}))
```

```
199 nan [[ nan nan nan]
[ nan nan nan]
[ nan nan nan]]
200 nan [[ nan nan nan]
[ nan nan nan]
[ nan nan nan]]
Prediction: [0 0 0]
Accuracy: 0.0
```

Not-A-Number

```
X = tf.placeholder("float", [None, 3])
Y = tf.placeholder("float", [None, 3])
```

LEARNING_RATE = 1e-10

```
W = tf.Variable(tf.random_normal([3, 3]))
b = tf.Variable(tf.random_normal([3]))
```

```
# tf.nn.softmax computes softmax activations
# softmax = exp(logits) / reduce_sum(exp(logits), dim)
hypothesis = tf.nn.softmax(tf.matmul(X, W) + b)
```

```
# Cross entropy cost/loss
cost = tf.reduce_mean(-tf.reduce_sum(Y * tf.log(hypothesis), axis=1))
# Try to change learning_rate to small numbers
optimizer = tf.train.GradientDescentOptimizer(
    learning_rate=LEARNING_RATE).minimize(cost)
```

```
# Correct prediction Test model
prediction = tf.argmax(hypothesis, 1)
is_correct = tf.equal(prediction, tf.argmax(Y, 1))
accuracy = tf.reduce_mean(tf.cast(is_correct, tf.float32))
```

```
# Launch graph
with tf.Session() as sess:
    # Initialize TensorFlow variables
    sess.run(tf.global_variables_initializer())
```

```
198 5.73203 [[ 0.80269563  0.67861295 -1.21728313]
 [-0.3051686  -0.3032113   1.50825703]
 [ 0.75722361 -0.7008909  -2.10820389]]
199 5.73203 [[ 0.80269563  0.67861295 -1.21728313]
 [-0.3051686  -0.3032113   1.50825703]
 [ 0.75722361 -0.7008909  -2.10820389]]
200 5.73203 [[ 0.80269563  0.67861295 -1.21728313]
 [-0.3051686  -0.3032113   1.50825703]
 [ 0.75722361 -0.7008909  -2.10820389]]
Prediction: [0 0 0]
Accuracy: 0.0
```


```
for step in range(201):
    cost_val, W_val, _ = sess.run(
        [cost, W, optimizer], feed_dict={X: x_data, Y: y_data})
    print(step, cost_val, W_val)
```

```
# predict
print("Prediction:", sess.run(prediction, feed_dict={X: x_test}))
# Calculate the accuracy
print("Accuracy: ", sess.run(accuracy, feed_dict={X: x_test, Y: y_test}))
```

전혀 학습이 되지 않음

12. Non-normalized inputs

```
xy = np.array([[828.659973, 833.450012, 908100, 828.349976, 831.659973],  
               [823.02002, 828.070007, 1828100, 821.655029, 828.070007],  
               [819.929993, 824.400024, 1438100, 818.97998, 824.159973],  
               [816, 820.958984, 1008100, 815.48999, 819.23999],  
               [819.359985, 823, 1188100, 818.469971, 818.97998],  
               [819, 823, 1198100, 816, 820.450012],  
               [811.700012, 815.25, 1098100, 809.780029, 813.669983],  
               [809.51001, 816.659973, 1398100, 804.539978, 809.559998]])
```



```

xy = ...
x_data = xy[:, 0:-1]
y_data = xy[:, [-1]]

# placeholders for a tensor that will be always fed.
X = tf.placeholder(tf.float32, shape=[None, 4])
Y = tf.placeholder(tf.float32, shape=[None, 1])

W = tf.Variable(tf.random_normal([4, 1]), name='weight')
b = tf.Variable(tf.random_normal([1]), name='bias')

# Hypothesis
hypothesis = tf.matmul(X, W) + b

# Simplified cost/loss function
cost = tf.reduce_mean(tf.square(hypothesis - Y))

# Minimize
optimizer = tf.train.GradientDescentOptimizer(learning_rate=1e-5)
train = optimizer.minimize(cost)

# Launch the graph in a session.
sess = tf.Session()
# Initializes global variables in the graph.
sess.run(tf.global_variables_initializer())

for step in range(101):
    cost_val, hy_val, _ = sess.run(
        [cost, hypothesis, train], feed_dict={X: x_data, Y: y_data})
    print(step, "Cost: ", cost_val, "WnPrediction:Wn", hy_val)

```

100 Cost: nan

Prediction:

[nan]

[nan]

[nan]

[nan]

[nan]

[nan]

[nan]

[nan]]

13.Normalized inputs (min-max scale)

```
xy = np.array([[828.659973, 833.450012, 908100, 828.349976, 831.659973],
               [823.02002, 828.070007, 1828100, 821.655029, 828.070007],
               [819.929993, 824.400024, 1438100, 818.97998, 824.159973],
               [816, 820.958984, 1008100, 815.48999, 819.23999],
               [819.359985, 823, 1188100, 818.469971, 818.97998],
               [819, 823, 1198100, 816, 820.450012],
               [811.700012, 815.25, 1098100, 809.780029, 813.669983],
               [809.51001, 816.659973, 1398100, 804.539978, 809.559998]]])
```

```
xy = MinMaxScaler(xy)
print(xy)
```

$$x'_j = \frac{x_j - x_{min}}{x_{max} - x_{min}}$$

```
[[ 0.99999999  0.99999999  0.         1.         1.         ]
 [ 0.70548491  0.70439552  1.         0.71881782  0.83755791]
 [ 0.54412549  0.50274824  0.57608696  0.606468   0.6606331 ]
 [ 0.33890353  0.31368023  0.10869565  0.45989134  0.43800918]
 [ 0.51436     0.42582389  0.30434783  0.58504805  0.42624401]
 [ 0.49556179  0.42582389  0.31521739  0.48131134  0.49276137]
 [ 0.11436064  0.         0.20652174  0.22007776  0.18597238]
 [ 0.         0.07747099  0.5326087   0.         0.         ]]
```

```
def MinMaxScaler(data):
    numerator = data - np.min(data, 0)
    denominator = np.max(data, 0) - np.min(data, 0)
    # noise term prevents the zero division
    return numerator / (denominator + 1e-7)
```

```

xy = ...
xy = MinMaxScaler(xy)
print(xy)

x_data = xy[:, 0:-1]
y_data = xy[:, [-1]]

# placeholders for a tensor that will be always fed.
X = tf.placeholder(tf.float32, shape=[None, 4])
Y = tf.placeholder(tf.float32, shape=[None, 1])

W = tf.Variable(tf.random_normal([4, 1]), name='weight')
b = tf.Variable(tf.random_normal([1]), name='bias')

# Hypothesis
hypothesis = tf.matmul(X, W) + b

# Simplified cost/loss function
cost = tf.reduce_mean(tf.square(hypothesis - Y))

# Minimize
optimizer = tf.train.GradientDescentOptimizer(learning_rate=1e-5)
train = optimizer.minimize(cost)

# Launch the graph in a session.
sess = tf.Session()
# Initializes global variables in the graph.
sess.run(tf.global_variables_initializer())

for step in range(101):
    cost_val, hy_val, _ = sess.run(
        [cost, hypothesis, train], feed_dict={X: x_data, Y: y_data})
    print(step, "Cost: ", cost_val, "WnPrediction:Wn", hy_val)

```

100 Cost: 0.152254

Prediction:

```

[[ 1.63450289]
 [ 0.06628087]
 [ 0.35014752]
 [ 0.67070574]
 [ 0.61131608]
 [ 0.61466062]
 [ 0.23175186]
 [-0.13716528]]

```